



PIP:

an integrated SDK for PicOS

version 0.5



1 Introduction

PIP is an attempt at an integrated SDK for PicOS. Before PIP, software development for PicOS involved a collection of command-line tools (mostly Tcl/Tk scripts), most notably `mkmk` and `picomp`, but also `genihex`, `piter`, and a few others, including manually invoked command-line utility programs, like `gdbproxy` and `gdb`, and independent FET programmers/loaders, like Elprotronic FET-Pro430-Lite. Underneath was `mspgcc` (the GNU tool chain for MSP430) and Cygwin's emulation of a UNIX environment (which we didn't have to worry about under Linux). Of course, there also was our “platform” consisting of PicOS, SIDE, and VUEE.

When contemplating PIP, I considered some alternatives for its development. At first sight, it would seem a natural choice to use Eclipse for its basis;¹ however, after a brief study, I decided against it. Here are the reasons:

1. Eclipse is a heavy-duty solution, whereas we would prefer something small. The size of its requisite luggage is intimidating. The sheer task of maintenance, e.g., keeping track of upgrades, would be very absorbing.
2. Eclipse is Java-based and, no matter how much you may want to disguise that fact, Java-oriented. A lot of its power looks much less impressive if Java is totally of no interest to the developer (as it happens in our case).
3. Software development under PicOS is not exactly a simple instance of the generic case of software development (in Java or C++). We have our own objectives and our own (idiosyncratic) targets, e.g., VUEE models (with their data files, including `udaemon`'s geometry data), multiple clones of the same uploadable Image differing by the node number, and our own organization of the “libraries” with the PicOS source code being a special (shared) part of all user projects.

Point 3 makes PIP sufficiently different from the “standard” case to cast serious doubts on the usefulness of Eclipse. All those idiosyncrasies of our setup would require special handling, meaning that on top of the other problems, the amount of work required to get something truly useful (as opposed to superficially flashy) within the Eclipse framework would most likely be quite nontrivial.

Consequently, I decided to hack instead something in Tcl/Tk, which has been our workhouse since day one. What we have now is already quite useful and effective. Especially for a newcomer it simplifies many things obviating the need to read our documentation, which is not always friendly. It should make the so-called quick start a bit quicker than it used to be.

2 Using X-Windows under Cygwin

While Linux would be (and should be) the preferred environment for PIP (and for the rest of our platform), the reality is that we still mostly use Windows + Cygwin. Within that framework, PIP is necessarily a kludge operating in the twilight zone between UNIX and Windows. Well, Cygwin is a kludge to begin with, and PIP merely pushes the idea to a new dimension. Its main window and the dialogs triggered by its menus operate exclusively in the Windows context, i.e., they formally need no Cygwin to live. Its text editor, however, needs X which runs under Cygwin. Although other X-Windows systems are available for Windows, the X server that comes with Cygwin is a natural choice, because PIP needs Cygwin (its set of utilities) for many other things as well.

In consequence of this kludgy setup, Windows windows must coexist with X-Windows windows within the context of something that we would like to call a single application. To make most sense out of this, you should run the X-Windows server in the *multiwindow* mode, i.e., without the background window covering the entire screen (which is known as the *windowed* mode). This is because, in the windowed mode, whenever you want to see an editor window, you will have to explicitly hide PIP's windows, and vice versa, which is going to be extremely annoying.

¹There exists an Eclipse-based SDK for `mspgcc` (see <http://homepage.hispeed.ch/py430/mspgcc/>). I cannot say anything about its quality, except that it is certainly not geared for PicOS or VUEE.



These days I find the multiwindow mode generally much better (friendlier) than the other mode, although for quite some time in the past I would stick to the windowed mode in my hopeless effort to conceal the fact that there's Windows underneath.

Assuming a full and reasonably up-to-date Cygwin installation, the recommended way to start the X server in the multiwindow mode is to double click **Cygwin-X → XWin Server** in the Windows startup menu (needless to say, you can create a convenient Desktop shortcut for it). The server reads a configuration file, `.XWinrc`, in your (Cygwin) home directory. If no such file is present, it behaves in some default way which is not bad. It shows an icon in the task bar, which you can right-click to open xterm windows (and possibly other stuff), as well as to terminate the server. There is no need to access Cygwin utilities in any other way, so you can forget about the (non-X) Cygwin terminal.

Here is my personal and trivial version of `.XWinrc` (which is not far from the default one):

```
# XWin Server Resource File (simplified by PG)
menu root {
    "Reload .XWinrc"      RELOAD
    Separator
    xterm                  exec    "xterm"
    Separator
}
RootMenu root
SilentExit
DEBUG "Done parsing the configuration file..."
```

As you can see, xterm client is practically the only application accessible from the menu. I have removed the indirection required to access it under the default setup, which you can find (for useful comments and reference) at `/etc/X11/system.XWinrc`.

With the X server operating in the multiwindow mode, most of PIP's kludgy nature will be hidden from your view. One little problem (which I couldn't do much about) is that the editor's windows cannot sensibly negotiate their way up through PIP's (Windows) windows, so they may open (partially) obscured. This is because the whole set of editor windows is considered second-class compared to Windows windows.

3 Invoking PIP

Just type:

```
pip
```

in an xterm. Note that it must be an xterm (or possibly some other X client), because PIP must know the identity of the X display. Strictly speaking, it is elvis, PIP's editor, who needs to know that identity, but elvis is invoked from PIP, so it boils down to the same thing.

In fact, this may be a bit confusing, because PIP will start without the X server (e.g., when invoked from a simple Cygwin terminal). It will even work and do things for you for as long as you don't try to edit a file, in which case PIP (or rather elvis) will fail complaining that it cannot talk to an X server.

You can go:

```
pip -D
```

(capital D), if you want to see debug messages produced by PIP on its standard output. Use this if you detect a (hopefully replicable) problem that you would like to report.

There is nothing wrong about multiple instances of PIP running at the same time for as long as they don't operate on the same project. There are no safeguards (locks) to prevent that, but messing things up this way requires a devious intention.

Generally, PIP tries to be permissive, quite likely beyond the safety zone. In particular, during some actions that take more time than just a click (like a compilation, for example), most of the action buttons will allow you to do other things in parallel with the action in progress. For as long as there is no interference, things will remain sane; however, it is not



impossible to mess them up. PIP does try to assess the sanity of data, often redundantly, to prevent crashes, but even if it doesn't fail formally, things may get confused. For example, changing the board configuration half way through a compilation is not recommended.

4 A walk through a simple project

We assume that everything has been set up already. A separate document, *Installation.pdf*, explains in detail what has to be installed. Go there first, if you are setting up things all by yourself.

We shall create a simple project from scratch and go through all the essential steps of its completion. For the praxis,² we shall use the “Hello World” program from Section 2.1 of *PicOS.pdf*. Even though this praxis already has an implementation (to be found in *Apps/EXAMPLES/INTRO/HELLO_WORLD*), we shall do it again. Besides, we will end up with something slightly bit more complicated than the old version.

4.1 Creating a new project

Start PIP, click **File → New project** and choose a directory. This will be a new directory, which should belong to the tree rooted at *PICOS/Apps*. The dialog that opens after the click is set by default at directory *PICOS/Apps*, which is the standard default container for project directories. Create a new directory there and name it, e.g., *HELLO_WORLD*.

Note 4.1.1: if you are on Windows/Cygwin, the dialog shows a **Make New Folder** button. On Linux, there is no button, but you can just type the directory name at the end of the path shown at the bottom of the dialog. The Windows dialog is confusing, so be careful. I may get around to programming a better (custom) version of this dialog some day.

Note 4.1.2: do not use file/directory names including spaces (PIP won't let you use them, anyway, but it cannot prevent you from creating ones spuriously under the standard open file dialog).³ They are considered illegal by PIP, because they confuse some of our underlying tools. This may be fixed at some point, but it isn't a high priority item.

When you select the directory and close the dialog, a new (project type) dialog will pop up: you have to decide whether the project is for a single program or for multiple programs. A multiple-program project is for those cases where the network needs nodes of different types, i.e., running different programs, which nonetheless belong to the same application. In a situation like this, you have to assign labels (alphanumeric keywords) identifying the different programs (see section 5 of *mkmk.pdf* if you are curious about details). As our simple project is going to comprise a single program, just press the **Single program** button and the dialog will go away.

The title of PIP's root window will now be showing something like “PIP 0.5, project *HELLO_WORLD*”. The left pane will present the initial list of the relevant files belonging to the project. For a new project, PIP puts there two files acting as stubs (to be filled out by you), but those stubs already constitute a compilable, loadable, and executable project (which does nothing, however). The two files are *app.cc* (the root source file) and *options.sys* (local re-definitions, if any, of the various parameters related to the system and to the node's hardware). This is the formally minimal set of files for any (single-program) praxis. Actually, the second file, *options.sys*, is optional, but it is a good habit to keep it around, even if it is empty.

Files in the left pane of PIP's main window (we shall call it the *tree view*) are organized into four groups named **Headers** (.h files), **Sources** (.cc files), **Options** (options files), **XMLData** (data files). Later, you will probably learn that there can be more than one *options.sys* file

²This is our name for a PicOS application.

³The way it works (under Windows/Cygwin) is that you first create a “New folder” (the name obviously contains a space in the middle) which you are then supposed to rename to something meaningful. So the dialog gives you the wrong hint.



(when the praxis is multi-program). The data files (their names mostly ending with .xml) pertain to the VUEE model of the praxis.

In addition to the above groups, which represent the project's files, i.e., ones stored in the project directory that you have just created (or its subdirectories, if any), the tree view may also show the directory of the board (i.e., the files describing the board) for which the project has been configured (as explained below). When the project consists of multiple programs destined for several different, the directories of all those boards may show up in the tree view.

4.2 Compiling for real hardware

The project can be compiled right away and even loaded into a board, although the latter is pointless, as the program literally does nothing. Nevertheless, let us try to compile it. When you click on the **Build** menu, you see that the only options available are for VUEE. This is because you haven't assigned a board to the project yet, so PIP doesn't know how to compile it for real hardware. To do that, click **Configuration → CPU+Board** which will produce a small "board selection" dialog. Don't touch the **MSP430** menu (this is our CPU type) and don't check the **Lib mode** box, but click on on the button under **Board** (initially showing something like ---) and select **WARSAW** (or whatever your board type may be). Then click **Done**. The tree view will now show (at the bottom of the previous set) a new entry (with a distinctive gray background) saying **<WARSAW>**. When you open it, you will see the list of standard headers describing the board of your project. If you look at the **Build** menu again, you will see that it offers two more choices appearing at the very top. They say **Pre build (WARSAW src)** and **Build (make)**.

To fully understand how a loadable image of your program is created, you should read *mkmk.pdf*. For now, suffice it to say that the operation is carried out in two steps: pre-building, which in a traditional setup corresponds to configuring the program for the specific environment, and the actual compilation. The outcome of the first step is a Makefile⁴ to be used in the second step.

So when you click **Pre build (WARSAW src)**, you will configure your project for compilation. You will see in the right (console) pane of PIP's main window messages produced by mkmk, which has been called to do the job. At the same time, the stripe above the tree view pane (the status line), which normally says **Idle**, will become **Running** and an advancing seconds counter will show up to the right. This means that some program is now running and (possibly) writing its output to the console. When mkmk has finished, the text **--DONE--** will appear at the end of the console's output and the status line will say **Idle** again.

Note 4.2.1: a program running in the console usually does not block PIP, in the sense that those features that can be sensibly used in parallel with that program are available (see the note in Section 3 on page 3).

Note 4.2.2: a program running in the console can be aborted by clicking the **Abort** entry from the **Build** menu (the very same action is also available from the **Execute** menu). The entry becomes enabled whenever the status line says **Running**.

The standard (and most flexible) way of compiling a PicOS praxis, we shall call it the *source mode*, involves the compilation of all system (PicOS) files needed by its program(s), including the kernel source code (see *mkmk.pdf* for details). This is because those files are heavily parametrized by various configuration constants whose purpose is to make sure that only the code fragments that are in fact needed are compiled in. This approach is OK, because the system is small enough to compile really fast. Of course, once compiled, the respective files will only be recompiled when their dependencies change, or when you "clean" the project, reconfigure it, or pre-build it, which action also removes all previously compiled object files.

An alternative way to compile a praxis, the *library mode*, involves an object library prepared in advance and (typically) associated with a specific board. The choice between the *source* and *library* modes is applicable to individual programs of the praxis; in a multiprogram case,

⁴For a multiple-program praxis, there may be more than one Makefile.



it may deal with multiple libraries prepared for the different boards configured with the different programs. The selection is determined by the **Lib mode** checkbox in the board selection dialog (mentioned at the beginning of this section). We need not concern ourselves with this issue right away.

Note that the **Build** menu shows two cleaning buttons: **Clean (full)** and **Clean (soft)**.⁵ The first button removes all files that can be recreated, including the Makefile(s) procured by pre-building. The second button works more like a traditional case of “make clean”, i.e., it removes the object files, but doesn't touch the Makefile(s). So the soft cleaning doesn't destroy a pre-build.

For our present exercise, the next step is to compile the praxis. When you click **Build (make)** from the **Build** menu, you will effectively execute make on the Makefile created in the previous step. The console will show the user output from this action, including any messages produced by the compiler and the linker. As before, the completion is signaled by the text --DONE-- appearing at the bottom of the console output; at the same time the status line reverts to **Idle**.

We did the two steps explicitly for education; however, you would have accomplished the same goal by clicking **Build (make)** directly. If no Makefile is present when a build is requested, PIP will carry out the pre-build action automatically, effectively emulating your clicking the Pre-build entry in the menu. Note, however, that a manual pre-build is always effective, in the sense that the full action of constructing the Makefile from scratch is carried out regardless of the status of any existing Makefile.

4.3 Editing files

Technically, the binary program produced in the previous step can be uploaded into the board. It doesn't make a lot of sense, however, because the program exhibits no activity that you could see. So we shall beef it up a bit before proceeding with that step.

Double click on the app.cc file (under **Sources**) in the tree view pane. This will open the file for editing in a separate window. Editing in PIP is handled by elvis⁶ which is a moderately popular VI-compatible, X-capable, and quite powerful text editor written (some time ago) by Steve Kirkendall and adapted by me to collaborate with PIP. It comes with extensive (built-in) documentation and is highly configurable. Most of the interesting configuration options of elvis are available from PIP's **Configuration** menu.

The stub in app.cc consists of a trivial root FSM with a single empty state. Edit it out to look to like this (feel free to be more creative):

```
#include "sysio.h"

fsm root {

    int led = 2;

    state INIT:
        ser_out (INIT, "Hello world\r\n");

    state ROTATE_LEDS:
        // Previous led goes off
        leds (led, 0);
        // Increment the led number modulo 3
        if (++led > 2)
            led = 0;
        // The led goes on
        leds (led, 1);

        // Wait for one second
```

⁵The selection for a multiprogram praxis is wider as the soft cleaning can be carried out on a per-program basis.

⁶Written with a small 'e' reserving the capital letter for the King.



```

        delay (1024, ROTATE_LEDS);
    }

```

If you pre-built the program in the previous step, there is no need to do it now, because, despite all that editing, you haven't changed anything that would affect the structure or configuration of system files needed by the program. For as long as you:

1. create no new files within the project
2. insert no new `#include` headers into the existing files
3. do not change the values of symbolic constants in `options.sys`
4. do not reconfigure the project (e.g., change the board)

there is no need to rebuild. In principle, the above situations could be detected by PIP/elvis and used to trigger pre-builds automatically, but I prefer to leave the decision to your discretion, the underlying philosophy being that giving power to experts makes more sense than comforting a handful of idiots.⁷ There is little penalty for pre-building "just in case", so in all those situations where you are not sure, just hit the button.

Regardless of what you do next, i.e., pre-build or build (compile), if you haven't saved the file before selecting the action in the menu, PIP will warn you that a file is being edited that has been modified but not saved and ask what to do. Your options are to save the file before continuing, continue with the old version of the file, or abort the operation. Needless to say, you can save the file from the editor window as well. Another thing that you may have noticed is that the tree-view name of a file being edited changes color to green or red (the latter means that the file has pending, unsaved modifications).

Close the editor window and try to build the program (press the second button on the **Build** menu). You will get an error message from the linker amounting to this:

```

/home/pawel/SOFTWARE/PICOS/Apps/HELLO_WORLD/app.cc:8: undefined
reference to `ser_out'

```

Double click on the error line in the console pane. The editor window (for `app.cc`) will pop up with the cursor pointing to the offending line. The problem is caused by the lack of a pertinent header file that would indicate to PIP (specifically to `mkmk`) that the library file containing `ser_out` should be compiled in. We shall address this problem in a little while.

4.4 C-tagging

Our program is extremely simple: it basically consists of a single file whose entire contents can be seen at once in a single (not so big) editor window. Nonetheless, we can use it to illustrate the search capabilities of PIP and elvis.

First let us complicate the program a bit by adding a second file (just to have a simple case of many). We shall encapsulate the code for rotating LEDs into a function. For that, let us create a new file named `leds.cc`. Right click anywhere within the **Sources** section of the tree view pane (e.g., on `app.cc`). In the menu that pops up, hit **New file** and name the file `leds.cc`. An editor window will open. Make the file look like this:

```

#include "sysio.h"

void rotate_leds () {

    static int led = 2;

    leds (led, 0);
    if (++led > 2)
        led = 0;
    leds (led, 1);
}

```

⁷One exception is reconfiguring the project which action forces **Clean (full)** thus automatically forcing a pre-build.



Then replace the corresponding piece in app.cc with a function call:

```
#include "sysio.h"

fsm root {

    state INIT:
        ser_out (INIT, "Hello world\r\n");

    state ROTATE_LEDS:
        rotate_leds ();
        delay (1024, ROTATE_LEDS);
}
```

Now close both editor windows, pre-build, and then build.⁸ This time you will get two errors (note that we didn't fix the previous one yet, so it is bound to show up again):

```
/home/pawel/SOFTWARE/PICOS/Apps/HELLO_WORLD/app.cc:9: undefined
reference to `rotate_leds'
/home/pawel/SOFTWARE/PICOS/Apps/HELLO_WORLD/app.cc:6: undefined
reference to `ser_out'
```

When you double click on the first of the two error lines, the editor window for app.cc will pop up pointing to the rotate_leds call. To fix this problem, you have to make sure that the second source file, the one you have just created, will also compile as part of the project. PicOS (or mkmk, to be exact), requires that those project files that should take part in pre-building and later compile into the project's Image be explicitly referenced from other files that are already marked for compilation (see Section 3.4 in mkmk.pdf). The only file compiled by default is app.cc. One way to make sure that leds.cc is compiled as well is to request it from app.cc by inserting (anywhere in the file) a comment looking like this:

```
//+++ leds.cc
```

In most cases, such special comments are put into header files and those header files are then included by the source files already belonging to the project, with app.cc acting as the root of the tree. In our simple case, we will just add the above line to app.cc and also insert there a header of rotate_leds, i.e.,

```
//+++ leds.cc
void rotate_leds ();
```

Suppose that you want to have a look at rotate_leds before fixing the problem (and your program is not as trivial as our educational example). Now, if you double click on rotate_leds in the editor window of app.cc, PIP will open the window for leds.cc with the cursor set on the definition of rotate_leds. This is called *C-tagging*. PIP builds (and at relevant stages rebuilds) its database of C-tags of all .cc and .h files belonging to your project (those listed in the tree view pane of PIP's main window) reflecting the locations in those files where things are defined, declared, and so on. Not everything is tagged this way, e.g., local (automatic) variables of functions and FSM states are not (assuming that those things are always easy to see in their usage context). Notably, system (i.e., PicOS) files can be tagged as well. For example, when you double click on one of the calls to leds in rotate_leds, PIP will open a (read-only) editor window with the system file sysio.h pointing you to the respective macro.

4.5 Access to system files

Click **Configuration → Options**. You will see a dialog defining three parameters associated with your project:

- the maximum number of lines stored in the console pane in PIP's main window (defaulting to 1000); any lines exceeding this limit will be discarded off the pane's top
- which PicOS files you want to see (and when) while navigating through the sources of your project

⁸Note that this time you have to pre-build, because you have added a new file to the project.



- which VUEE files (i.e., ones belonging to the VUEE system model) you want to see

Note that VUEE is a separate “system” with its own source files implementing SIDE models of the PicOS stuff and the network environment. If you are running your program virtually under VUEE, you may be interested in seeing the interiors of that model, rather than the interiors of PicOS.

The options regarding your access to system files (defined separately for PicOS and VUEE) are:

Never → the system files are never C-tagged and never searched by PIP (with the search engine described in Section 4.6 below).

Tags, R/O → the system files are C-tagged, but not searched. When referenced by C-tags, those files will open read-only (so you won't be able to modify them inadvertently).

Always, R/O → the system files are C-tagged and also searched. They always open read-only.

Always, R/W → the system files are C-tagged, searched, and always open read-write.

The last option is for PicOS development (from inside a project). Note that the “protection” offered by PIP's configuration options is purely advisory. You can always edit a system file from outside PIP and do with it whatever you please. Also, a read-only file can be overwritten by elvis with the `w!` command, as long as you have formal write access to the file within the file system.

The default settings are **Tags, R/O** for PicOS and **Never** for VUEE. Note that they apply on a per-project basis, unlike, e.g., color schemes for elvis, which are stored globally (for the user).

Note 4.5.1: if you change the setting for PicOS system files to **Never**, the board directory entry (or board entries, for a multiprogram project) will disappear from the tree view. Formally, board descriptions belong to the system, so if you declare that you do not want to view system files at all, PIP will hide the board definition(s) from your view.

The C-tags database for your project's files is updated whenever any of those files is modified (or created), but not until the file is saved by elvis. For the system files, the situation is a bit more complicated (and is different for PicOS and for VUEE).

For PicOS, only those system files that your project actually needs are C-tagged. Their set can only be known after you have pre-built the project. Consequently, no PicOS files are accessible via C-tags clicks if the project hasn't been pre-built yet, or after you have fully cleaned it (**Build** → **Clean (full)**). In the more complicated case of a project consisting of multiple programs, the list of PicOS files to be C-tagged is determined from all pre-builds combined (by inspecting the current collection of the project's Makefiles).⁹

For VUEE, assuming the C-tags option is on, the C-tags are computed once for all, for the complete set of VUEE files. These are the files in the VUEE/PICOS directory (of the VUEE package).

4.6 Searching

We are now ready to try to attend to the first problem, i.e., the missing header file for `ser_out`. A simple idea to figure out what that file might be would be to double click on the respective function call in `app.cc`. But nothing happens when you do that, except for the line showing up in the console pane to tell you: “Tag `ser_out` not found”. Why not? Because, as explained in the previous section, to be C-tagged, a PicOS file must be referenced from the project, and this is exactly our problem: the file is not referenced!

⁹It is slightly more complicated (technically, not semantically) when the library mode is used. In that case, the list of system files is determined from the library content.



In a situation like this you may want to resort to searching. Click **File → Search** to bring up the search window. Select **WD** and check **Case**. The first widget determines the interpretation of the search string (as a word, or a sequence of words), the second says that the case of letters matters. You may select **Only** from the **Sys** list to indicate that you want to search *only* system files. Then type `ser_out` into the **String** field. Alternatively, you can double click again on the function name in `app.cc`. When the search window is open and the tag is not found, it is automatically inserted as the search string. Now press the **Search** button.

You will see a lot of matches looking mostly completely uninteresting (lots of `params.sys` files in board definitions). We can easily narrow down the search to what we really care about, i.e., the header files. For that, enter `h$` into the **FN Pat** field. This is a regular expression to be matched against the file name. Clear the output area with the **Clean** button and try again. This time it looks much better. There is less than a handful of matches, and you can clearly see that the requisite header file is:

```
.../PIC0S/Pic0S/Libs/Lib/ser.h
```

so you have to insert:

```
#include "ser.h"
```

somewhere at the beginning of `app.cc`.

Every match presented in the output pane of the search window begins with a highlighted header listing the file name and the line number in the same format as error locations produced by the compiler. When you double-click on such a header, and the current option settings (Section 4.5) allow you to view the file, an editor window will pop up with the file positioned at the respective line. Every match is shown in the search window as a range of lines whose width is determined by the **Bracket** field from the lower line of widgets. The default setting of 5 means that the range will include 5 lines in addition to the matching one (i.e., 6 lines altogether), with the matching line located approximately in the middle. The matched string is highlighted within the line.

You could narrow down the search even more by adding an opening parenthesis after `ser_out` in the **String** field. When you try it (together with `h$` in **FN Pat**), you will get exactly one match. With word matching (selected by **WD**) it doesn't matter how many spaces appear between the function name and the parenthesis; what does matter instead is the exact succession of *words* with any blanks between them ignored. The three options for the interpretation of the search string are:

RE → regular expression¹⁰

ST → string, i.e., simple string matching without interpreting any characters as special

WD → a sequence of words, possibly separated by blanks (but no newlines)

For the last option, a word is either a sequence of consecutive letters and/or digits and/or the underscore character, or any other character standing by itself for a complete word. Thus, for example:

```
words @$ 11@( _ABC1/
```

stands for 9 words:

```
words, @, #, $, 11, @, (, _ABC1, /
```

which will match any of these:

```
words  @#  $11@(_ABC1/
words@  # $ 11 @ ( _ABC1 /
```

but not:

¹⁰See http://www.tcl.tk/man/tcl8.5/TclCmd/re_syntax.htm for a detailed description of regular expressions in Tcl.



```
words@#a$11@(_ABC1/
words @ # $ 11 @ ( _ABCD1 /
```

Whether the case of letters is relevant or not depends on the **Case** option.

The **Sys** selection allows you to choose whether you want to search through the PicOS system files, with the **VUEE** checkbox optionally throwing in the VUEE files to the system pool. The options are:

None → only the project files will be searched (no system files at all)

Proj → only the system files associated with the project (based on the current pre-build) + all VUEE files, if the VUEE box has been checked, will be searched

All → all system files will be searched

Only → *only* all system files will be searched, i.e., not the project files

Note that in the second case the project must have been pre-built (see Section 4.5) for the PicOS files to be searchable (only then can PIP know the exact pool of files to search).¹¹ This does not apply to the VUEE files (if the **VUEE** box is checked), which are always searched *all*.¹²

The lower widget line of the search window contains action buttons and fields. The role of those fields is to limit things. For example, the already described **Bracket** field limits the width of the displayed neighborhood of a matched line. **Max lines** determines the maximum number of lines stored in the output area (defaulting to 1000). **Max cases** limits the number of matches (the default limit is 256); searching stops when that number is reached.

By putting a (Tcl) regexp into the **FN Pat** field you restrict the population of files to those whose names match the specified pattern. This field is typically used to limit the search to source or header files. By checking the **!** box in front of the field, you negate the interpretation of the pattern, i.e., it marks the files that *will not* be searched.

The colors used to highlight the headers of the matched sequences presented in the output pane, as well as the matching fragments within those sequences, are configurable with the **H** and **M** buttons. Just press them and see what happens.

The meaning of the action buttons **Search**, **Clean**, and **Close** is rather obvious. The remaining two buttons, **Edit** and **New**, allow you to edit and create any file from PIP, not only a file related to the project, which may be useful when developing PicOS or VUEE.

4.7 Uploading programs into boards

Let us fix our program, i.e., make sure that these lines have been added to app.cc:

```
#include "ser.h"
//+++ leds.cc
void rotate_leds ();
```

Now we have to pre-build, because we have added a new header file to the project. When we do that and subsequently build the program, there will be no more errors. The program is ready to be loaded into the board.

For that exercise, we shall assume that your hardware connection to the board consists of the FTDI TTL232R USB UART dongle and the MSP-FET430UIF programmer, both devices connected to the PC via USB cables. All the requisite drivers are present as per Installation.pdf.

Connect the two devices and make sure that the board is powered up. Click **Execute → Start piter** to open a terminal window that will allow you to communicate with the board via its UART, once the board gets to running the program.¹³ Select the respective COM device of

¹¹For a program built in the library mode, no pre-build is necessary, because the set of PicOS files needed by the program is determined based on the library content.

¹²This reflects the fact that all those files are compiled into any VUEE model.

¹³The program is documented in Serial.pdf.



the FTDI dongle. If you don't know it, hit the **Scan** button; then the device menu will only show those COM devices that respond. The dongle is usually easy to spot as something new and different from what you are used to. For subsequent sessions, the same dongle will tend to map to the same COM device.¹⁴ Select the speed of **9600** bps and **Direct** protocol, also make sure that the **Bin** box is not checked (these should be the defaults, anyway). When you hit **Save & Proceed**, your settings will be saved inside the project, so the next time around you will see them as the new defaults.

With MSP-FET430UIF, and under Cygwin, you have two options regarding the loader program.¹⁵ In PIP's main window, click **Configuration → Loaders**. You can choose between Elprotronic FET-Pro430-Lite and msp430-gdb (this is actually a debugger also providing loader functionality) communicating with the board via msp430-gdbproxy. Let us explore the second option. Select **Use** for msp430-gdb and make sure that **FET430UIF** is shown in the menu as the programmer device. Then click **Done**.

Now click **Execute → Upload image**. Note that this item only becomes enabled (clickable) after a successful compilation (an Image file is present in the project's directory). A window will pop up with two panes. The left pane shows the output of msp430-gdbproxy, while the right pane is for msp430-gdb, which is initially not active. If the bottom of the left pane says something like PROXY CONNECTION ESTABLISHED, it means that msp430-gdbproxy has successfully connected to the device and is ready to receive commands from msp430-gdb. Otherwise, something is wrong and, usually, some hints regarding the reason are displayed in the pane.

One possible reason for failure at this stage is the wrong firmware in your programmer, which is quite likely, if you have never before used the device with msp430-gdbproxy. In such a case a pertinent message should show up in the left pane. Click the **Update FET** button and wait for a few seconds. If everything works as desired, you should see a bunch of messages in the left pane indicating progress in feeding new firmware into the programmer. Wait until you see REPROGRAMMING COMPLETE and then hit the **Restart** button.

Note 4.7.1: never try to reprogram Olimex's JTAG TINY programmer. When you do that, you will likely damage the device.

Once the proxy has connected to the board, the next step is to choose the Image file to upload. There is not much choice in our simple case (there is only one Image file), but sometimes (e.g., for multiprogram projects) this step may be nontrivial. The menu at the bottom of the upper right column of buttons lists all Image files that PIP sees as possible candidates for upload. The only file present there in our case is named Image. This is the ELF¹⁶ output produced by the linker at the end of build (compilation).

Push the **Connect** button to start msp430-gdb for the (selected) Image file. The right pane should come to life displaying the initial output from msp430-gdb. If everything works fine, you should see the text GDB CONNECTED TO DEVICE at the bottom. Also, the three buttons under **Connect** should become enabled. Click **Load**. After a few seconds the left pane should show a bunch of “memory write” messages from msp430-gdbproxy. Eventually you should see the text PROGRAM LOADED SUCCESSFULLY at the bottom of the right pane. The program has been loaded. Click **Run** to reset the board and start the program. In piter's window you should see a standard PicOS header (written to the UART after every reset), followed by “Hello world”. The three LEDs will begin their rotation switching every second. You can now disconnect the board from the programmer, but you may want to close the loader window first, so msp430-gdbproxy will not feel offended.

Note that the right pane of the loader window gives you full (manual) access to gdb with the three buttons (**Load**, **Run**, **Stop**) providing merely shortcuts for the most typical operations.

¹⁴This isn't true on Linux where the numbering of ttyUSBx devices reflects the order of their connection.

¹⁵On Linux, msp430-gdb + msp430-gdbproxy is your only option. This may change.

¹⁶Executable and Linkable Format. This is the standard format of object modules and executables produced by the GNU linker.



The input field at the bottom can be used to enter directly any command. In particular, you can carry out the (previous) load procedure with these commands:

```
monitor erase all
load Image
continue
```

You can also set breakpoints and debug the program. The document by Steve Underwood, known as mspgcc.pdf in our set, provides a few useful hints.

4.8 Building and running a VUEE model

Building a VUEE model of our simple application is quite easy and involves no additional programming. One extra thing we have to worry about is a data file to the model, which describes such parameters as the population of nodes, their hardware configuration, their spatial distribution, the properties of the radio channel, and a few more items. There is no network in our case (we haven't programmed any radio activity into the node), so we can get away by borrowing the trivial data set from the old HELLO_WORLD praxis that comes with the package.

In the tree view pane, select (left click) **XMLData** header and then right click on it. From the menu that pops up select **Copy from**. Alternatively, having selected **XMLData**, you can click **File → Copy from** for the same effect.

Navigate to **EXAMPLES/INTRO/HELLO_WORLD/data.xml** and “open” that file. It will be copied to your project and appear under **XMLData** (as data.xml). When you look at it, you will see that it describes a trivial “network” consisting of one node. Now we are ready to run the program through VUEE.

Click **Configuration → VUEE** to configure the model. Check **Always run with udaemon**, and select data.xml for **Praxis data file**. No need to touch the remaining widgets. Click **Done**.

Click **Build → VUEE** to compile the model. When you see **--DONE--** in the console pane (and **Idle** in the status line), the model is ready. To run it, click **Execute → Run VUEE**.

In udaemon's window that will pop up shortly, type 0 into the **Node Id** field (there is only one node numbered 0 in our “network”) and click **Connect** to connect to the node's UART. You will see the “Hello world” message, as before, except that now it is coming from a virtual board. To see the LEDs, select **LEDS** from the menu in udaemon's window and click **Connect** again, this time to connect to the node's LEDS module.

It is all much more impressive with a geometry and window preset data file for udaemon causing the model's complete visualization to pop up at the click of a single button. Read **VUEE.pdf** for details.

Note that VUEE models can be disabled for a project by checking a box in the **Configuration → VUEE** dialog. This has the effect of dimming/disabling the VUEE-related items in the **Build** and **Execute** menus and makes sense for those projects that are not meant to be VUEE-compatible, e.g., hardware tests.

4.9 Running external programs from PIP

As a crude generic solution for problems that are not directly addressable by PIP's widgets, there is a way to execute any (external) command in the project's directory. The bottom line of the console, normally disabled, provides a means to enter input to such programs. Their output will appear in the standard output area of the console.

One special case of an external program is gdb used to debug a VUEE model. Note that debugging a program running in real hardware can be done from msp430-gdb (used as the loader), although personally I have practically never resorted to that method. There is something about programming in PicOS that makes debugger-assisted program analysis considerably less appealing in comparison to other systems. In the vast majority of cases the problem is either conceptual (and the debugger cannot possibly help), or it can be fixed much faster by a simple inspection of code. Even for VUEE, most of the situations where the debugger does help involve errors in the VUEE emulator rather than in your praxis code (so



it is more for VUEE development than anything else). Then, the debugger's role usually reduces to pointing at the line that triggers the segmentation violation, or whatever the exception happens to be.

To see how the debugger works, start by compiling the VUEE model with debugging enabled. Select **Build → VUEE (debug)**; this will add the -g flag to the options of the SIDE compiler causing debug information to be appended to the model's executable. After the compilation successfully completes, hit **Execute → Run VUEE (debug)**. PIP will start gdb in the console window giving it the SIDE executable file as the argument. The bottom (input) line will activate, and you will be able to enter commands to the debugger. Note that regardless of the VUEE configuration settings (in **Configuration → VUEE**) `udaemon` is not automatically started in this case. You can always start it manually (**Execute → Run `udaemon`**) if needed.

Enter:

```
run data.xml
```

into the console's input line to start the model. You may want to enter:

```
run data.xml +
```

instead. The second argument (standing for the output file for the VUEE model) indicates that the output should be written to the standard error (rather than the default standard output) with the net effect of flushing the output on every line (so you can see it right away). This problem is specific to Cygwin/Tcl8.5 and results from one of the never ending list of obnoxious mismatches. Another problem is the Control-C doesn't seem to work for gdb under Cygwin (I distinctly remember that it used to work fine, or maybe it still does on every second day?).

You can abort any program running in the console with **Execute → Abort** (or **Build → Abort**, both entries trigger exactly the same action).

When you select **Execute → Run program**, you will be shown a simple dialog to enter the command you want to execute in the console window. This is potentially dangerous, so you should know what you are doing. In particular, the program can be `sh`. In such a case, however, a better idea might be to click **Execute → XTerm** and simply open an xterm client in the project's directory. The xterm is completely detached from PIP; in particular, it won't go away when PIP exits.

5 Elvis configuration

Elvis can be configured in many ways and in many aspects, and it is certainly beyond the scope of this note to discuss them all. The editor is equipped with an elaborate help file which can be easily accessed from the program itself. From the viewpoint of PIP, one useful possibility is to assign different color schemes to different file types, which may help in editing/viewing lots of files at once.¹⁷

By default, PIP provides elvis with a single (default) color scheme to be used for all files. Whatever other schemes you define, they will be stored in PIP's global configuration file, `.piprc` in your home directory, and applied globally to all your projects.

To define a new elvis scheme you must be inside a project. This is because the Configuration menu only works if a project is currently open. Click **Configuration → Editor schemes**. A dialog will pop up. Click **New** to open another small dialog. Name your scheme (any alphanumeric string will do) and indicate which of the existing schemes will be used for its basis. Initially, there is nothing to choose from except for the default scheme. Click **Create** to proceed.

In the configuration window that opens next you can assign colors to the various elements (faces) displayed in elvis's windows (you will find their description in the editor's help pages). The first entry, labeled *normal*, is most important: its base color is simply the color used to display (general) text, while the background color applies to the window background. The

¹⁷Note that references to C-tags may bring in lots of elvis windows.



alternate color stands for the second choice for the base color to be used when (according to elvis) it will give a better contrast. Feel free to experiment.

Note that the configuration of a face may be “linked” to some other face. That means that the definition is the same as in the linked-to face, except for the items that are explicitly marked as different in the linking one. The boxes labeled B, I, U, O, stand for bold, italic, underline, and boxed, respectively, and refer to text attributes.

In addition to colors, you can also select the font size for text and, if you know what you are doing, enter explicit, textual, extra configuration commands to be executed by the editor when it starts.

Having completed the definition of a new scheme, hit **Done**. The scheme will now be available in the menu in the left upper corner of the left (schemes) pane in the previous dialog (which should remain open). You can assign your schemes to file types using the set of widgets in the right (assignment) pane of the schemes dialog. Those widgets are organized into rows, with one row assigning one scheme to some subset of files. The way this is done should be intuitively clear. You select the file type (note that those types correspond to the tree view headers) and where it is coming from (project, system, etc.), with the truly tricky cases resolvable by explicit regular expressions applied to file names. Whenever an editor window is to be opened for some file, the rows will be scanned from top to the bottom, and the scheme from the first matching row will be used for the file. If no match is found, the default scheme (which is always defined) will be assumed.

6 Multiprogram projects

Section 5 in mkmk.pdf describes how multiprogram praxes are handled by the command-line tools. Here we show how it all looks from PIP through an easy to follow example.

Start by inspecting the project at Apps/EXAMPLES/SIMPLE. This is a single-program project (there's just a single app.cc file in the directory) illustrating how to implement simple one-hop RF communication. For a real-life demonstration do this:

1. Build the project, e.g., for the WARSAW board.
2. Upload the Image file into two (WARSAW) nodes.
3. Interface both nodes to COM ports, e.g., via separate FTDI dongles. Start two instances of piter at 9600 bps, mode Direct¹⁸ and make sure they connect to the proper COM ports.
4. Whenever you enter:

s ... some text ...

on one node, you will see the text echoed by the other node.

You can carry out the same exercise virtually in VUEE. For that:

1. Build the VUEE model of the praxis as described in Section 4.8.
2. Run the model with udaemon using data.xml for the VUEE data file.
3. In udaemon, open the UARTs for nodes 0 and 1.
4. Do the same as in point 4 of the previous exercise.

For the multiprogram demo we have to come up with at least two different programs to be run by different nodes within the framework of the same praxis. To this end, Apps/EXAMPLES/SIMPLE_TWO contains a derivative of the SIMPLE project with the sender/receiver functionality split between two separate node types, such that a single node can be either a sender or a receiver, but not both.

6.1 Program labels

When you open the SIMPLE_TWO project, you will see under the **Sources** heading in the tree view two files: app_snd.cc and app_rcv.cc constituting the *roots* of the two programs.

¹⁸ You can run multiple instances of piter from the same PIP session (Execute → Start piter).



The parts “snd” and “rcv” of the file names are called *labels*. They play the role of identifiers of the multiple (two in this case) programs.

In the case of SIMPLE_TWO, the two root files is all there is, i.e., there are no more source files contributing to the project. This of course would be different in a more serious case. The way the source files are attributed to the respective programs is described in Section 4.4 (and also in Section 3.4 of mkmk.pdf). Note that the programs may share some files.

When you create a new project, you can choose the multiprogram option in the **Project type** dialog (see Section 4.1) in which case you will have to provide the labels for the multiple programs. The number of different labels you enter in the **Project type** dialog will determine the number of programs in the project, and PIP will create that many stub `app_XXX.cc` files, with XXX replaced with your labels. A label can be any alphanumeric string (letters and digits) starting with a letter.

Note that to be formally deemed “multiprogram”, in the sense of employing labels, a project doesn't absolutely have to consist of multiple programs. For example, try creating a new project (say a new version of HELLO_WORLD) entering a single label in the **Project type** dialog and clicking the **Multiple programs** button. Basically, everything will proceed exactly as before, except that your root (app) file will be labeled, i.e., its name will look like `app_XXX.cc` where XXX stands for the label you have entered into the **Project type** dialog. So the simplest case of “multiple” is in fact “one” in this case. While labeling the single program in a single-program project is allowed, it brings nothing new to what we have already learned about single-program projects. In projects actually consisting of multiple programs, the labeling is necessary as a way of telling the programs apart. From now on, we shall assume that we in fact deal with multiple programs (“multiple” means at least “two”).

6.2 Build options for multiprogram projects

Open the SIMPLE_TWO project and click **Configuration → CPU + Board**. You will see that the **Board selection** dialog still includes the **Multiple** checkbox, even though there's no question as to the multiple-program status of the project. This box does make sense. Its one purpose is to tell whether the multiple programs require different boards and also offering a certain simplification if they don't.

When you uncheck the **Multiple** box, you basically say that all programs of the project are for the same board and they can be compiled as such. There will be only one board to choose in the **Board selection** dialog, and the choice will apply to all programs. This also implies that all programs share the same local configuration options which, at the project level (as opposed to the board level), are described in the single `options.sys` file in the project's directory. In the **Build** menu you will see a single pair of pre-build/build actions applicable to the entire project: the respective action takes care of all programs at the same time (they are all pre-built and subsequently built together). In more technical terms, it means that there is a single Makefile for all programs.

When you successfully build the project, you will see that the outcome of that action amounts to two Image files named `Image_snd` and `Image_rcv`, i.e., the files are tagged with the labels of their programs.

Note 6.2.1: the simplest shortcut to see the full contents of the project's directory is to open an xterm (**Execute → XTerm**) and executing `ls`.

Note 6.2.2: each Image file occurs in two versions (this applies to all projects, not only the multiprogram ones): one without an extension (this is the ELF file mentioned on page 12) and the other with extension `a43`. The latter is the Intel HEX format of the ELF Image, which is used by some loaders, e.g., the Elprotronic FET-Pro430-Lite. If you decide to upload the programs into real boards, you will see the (respective) Image files available as options to the (respective) loaders.

Now click **Configuration → CPU + Board** again and check the **Multiple** box. The shape of the dialog will change allowing you to specify multiple boards on a per program basis. The project will be built differently now: the **Build** menu lists separate pre-build and build actions for each program. Although composite actions (**Pre-build all** and **Build all**) are available, they



are merely shortcuts for sequences of independent actions carried out individually for each program. Technically, this means that each program gets its own Makefile which, as one can easily guess, is tagged with the program's label.

In the Multiple build mode, different programs can specify different local configuration options, which you provide in separate labeled options files, e.g., `options_snd.sys` and `options_rcv.sys` for the project at hand. Those files are not mandatory; if absent, PIP will try to use the global `options.sys` file for all (both) programs. That file is not mandatory, either: if there's no applicable options file, PIP (or rather mkmk to be exact) will assume that there are no local options. Local options files can be considered a legacy approach to specifying the attributes of the praxis's target hardware, which role has been obsoleted and obviated with the introduction of board descriptions. Unfortunately, we still use them (by tradition) for adjusting or correcting some hardware parameters of the board(s), like the UART rate, the presence or absence of a particular device, etc., while their role should be rightfully confined to specifying configuration parameters directly affecting the praxis programs. In addition to clearly separating concerns, a consistent adherence to this principle will also facilitate board-specific libraries (see Section 7).

Note 6.2.3: one remaining well-defined role of the global `options.sys` file is to define the configuration of VUEE components (see Section 3.4 of `VUE_doc.pdf`). Note that VUEE is only concerned with the global `options.sys` file and always ignores the labeled versions.

Even when the **Multiple** box is checked in the **Board selection** dialog, you can still select the same board for both (all) programs of the project. This will not affect the independent status of the programs and their independent builds. To sum up:

1. If all programs of a project are for the same board, you have an option of treating them separately or combining them into a single build.
2. If at least some programs require different boards, then you must check the **Multiple box** in the **Board selection** dialog, which will result in independent builds for all programs.

Regardless of whether the Multiple box is checked or not, the Image files built in a multiprogram project are always program-specific and they are always tagged with the respective program labels.

6.3 VUEE models of multiprogram praxes

Most of what matters for this leg of the project, we have already said in the previous section. The details, including some of the issues that must be resolved (mostly by PiComp) when turning the set of programs comprising a praxis into a VUEE model, are described in `PiComp.pdf`. One thing you should remember is that the program will use `options.sys` to determine the configuration of model components. For example, in `SIMPLE_TWO`, the file contains a single line looking like this:

```
#define VUEE_LIB_PLUG_NULL
```

which declares that the model needs the NULL plugin (see Section 3.4 of `VUE_doc.pdf`).

Make sure that the VUEE model has been configured, i.e., click **Configuration → VUEE**, check the **Always run with udaemon** box, and select `data.xml` for **Praxis data file** (no udaemon geometry file is provided in the project's directory, but the model can live without it). Then click **Build → VUEE**, wait until the compilation completes, and click **Execute → VUEE**. By looking at `data.xml` you can tell that node 0 is the sender (snd) and the other node is the receiver (rcv). Open the UARTs of the two nodes. When you type:

```
s ... some text ...
```

into the UART of node 0, you will see the text echoed by node 1.

7 The library mode

In Section 4.2 we mentioned two build modes for a project, or rather a project's program, because the multiple programs of a multiprogram project can be build in mixed modes, as



long as their builds are in fact independent (i.e., the **Multiple** box in the **Board selection** menu is checked – see Section 6.2). The library mode is selected by checking another box in the **Board selection** menu, the one that appears under the board name for the given program. If the **Multiple** box is not checked (and all programs compile as a single build), the library/source mode is selected globally for the entire project.

A library is a precompiled collection of object files. The set of system options for such a precompilation, as well as the exact collection of files contributing to the library (i.e., the configuration of system modules) are determined by the board definition. This means that libraries are naturally (and rather strongly) associated with boards.

By invoking mkmk manually (see Section 4 in mkmk.pdf for details), one can create libraries that can be additionally parameterized by local options (supplementing or modifying those associated with the definitions of the corresponding boards). Within the framework of PIP you can only create and subsequently use only “pure” libraries, i.e., ones whose shape is solely determined by the configuration of their boards. As mentioned in Section 6.2, the whole idea of augmenting (or correcting) board definitions by local options is probably flawed and should be avoided in serious (production) projects.

Let us return to SIMPLE_TWO for an illustration. We shall reconfigure the project to build in the library mode. Open the Board selection dialog (**Configuration** → **CPU + Board**); it doesn't make much difference whether **Multiple** is checked or not, but let us keep it unchecked for now (to make things a little bit simpler). Later you can try to go through basically the same drill with multiple boards. Check the **Lib mode** box under the board name (which should be saying **WARSAW**).

With **Multiple** unchecked, both programs of the project will build together (compiling for the same board). The pre-build action from the **Build** menu now says **Pre-build (WARSAW lib)**; click it. Unless you have already created a library for the WARSAW board, PIP will complain telling you that no library is available for WARSAW. You can build such a library by clicking **Configuration** → **Create lib for WARSAW**. Do that. You will see a bunch of messages in the console pane reminiscent of a two-stage build of a regular program: there will be a pre-build stage followed by a compilation. Finally, assuming that the operation is successful, the console pane will say **-LIBRARY CREATED-**, and the status line will revert to idle. The library has been created. PIP stores libraries in directory **LIBRARIES** of the PICOS tree. This directory is set up on the first library build, if it doesn't exist. The libraries are named after their boards; thus, the operation described above will yield subdirectory **WARSAW** in **LIBRARIES** containing the library for the WARSAW board.

The **Configuration** menu includes actions for building libraries for all the boards configured into the project.¹⁹ Once built, a library can be rebuilt (just by clicking the respective item in **Configuration**), but PIP will warn you (and ask for confirmation) if a library for the given board already exists. No dependencies are automatically detected. Normally, libraries have to be rebuilt only when some of the system files they incorporate (directly or indirectly) have been modified.

Our next step will be to actually build the project. Now the pre-build action will succeed. Note that both pre-building and building are now much faster than in the source mode, because much of the effort has been already invested into the library.

You should remember that, in a library build, any options files present in the project's directory (and sought according to the rules outlined in Section 6.2) have no effect on system parameters, although they can still affect the project programs. And, of course, the VUEE model (which doesn't distinguish between the library and source build modes) still interprets options.sys for the requisite components.

A board description must adhere to certain rules to allow for its library to be created. Those rules are described in Section 4 of mkmk.pdf. In a nutshell, the file params.sys in the board's directory must specify the list of files to be included in the library.

¹⁹If the project uses multiple boards, there will be also a compound action to build them all one by one.



8 Comments

People used to accomplishing their goals solely by clicking buttons will be disappointed by the fact that PIP leaves considerably more to textual input than “more professional” solutions. For example, other SDKs would let you add files to the build by clicking on them rather than forcing you to insert special comments into other files (see page 8). Personally, I find it more natural to reflect the fact that something is needed by my program within the specific fragment of code that depends on the requisite item than to rely on a mark in some obscure internal description of the project. Traditionally, such information has been handled quite well textually through header files, and it is still difficult to beat its “expressive power” by clicking buttons.

Another thing that PIP does little about is hardware description. Traditionally, we have been storing hardware (board) descriptions in a special directory inside the PICOS tree (see PICOS/PicOS/MSP430/BOARDS) as collections of header files. While it is conceivable to visualize those descriptions better, e.g., turning them into a bunch of selections, entry boxes, radio buttons, and so on, I wouldn't want to loose in the process the convenience of being able to keep there comments (sometimes quite elaborate) as well as miscellaneous pieces of code that fit there better than anywhere else. So I still believe that nothing beats the simplicity of being able to mix so much “multimodal” information in a short file that you can see all at once in an editor window.

One thing that probably would be useful is a special collection of simple widgets for editing and navigating board descriptions. You can do it now by using the buttons in the search window, but perhaps something specifically addressing this particular end would not be completely out of place.

