



Accessing Olsonet git Repositories



August 11, 2010

© Copyright 2010, Olsonet Communications Corporation.
All Rights Reserved.

Introduction

We have given up the CVS repository at cs.ualberta.ca. The primary reason was a security breach affecting some departmental systems and causing a prolonged disconnection of our server from the world, which left us with no easy means to synchronize our changes at a rather critical moment (as it usually happens). Using this incident as a pretext, we decided to switch to *git*. The conversion of the original CVS repository to git (carried out on June 28, 2010) was pretty much automatic (git-cvsmimport). The single original repository (holding three packages: PICOS, SIDE, and VUEE) has been split into three separate git repositories. The original structure of commits, branches, and tags has been preserved and, as far as I can tell, all the different snapshots/branches that were accessible under CVS are still easily retrievable via git.

This document explains how to access the present repositories of Olsonet hosted at <http://unfuddle.com>. This is an updated version of the note that I prepared on June 28 for the then free subscription to unfuddle, which was considered temporary. There are very few formal changes with respect to that note. The most obvious difference is a new shape of the access URLs.

Please become acquainted with the basics of git before reading the rest of this document. While I give you some examples illustrating typical commands and scenarios, they shouldn't be used as a replacement for elementary understanding of what you are doing. Here are two useful links:

<http://progit.org/book/>
<http://www.kernel.org/pub/software/scm/git/docs/user-manual.html>

The first one is a friendly, reasonably short, and quickly absorbed online book on git; the second is the more or less official manual (not quite complete, but very useful if you have to find out about details of some of the more tricky features). And, of course, you can always resort to the (ugly as usual) man pages.

Local prerequisites

Make sure that you have a decent version of git installed on your system. I have done my work so far under Cygwin (git comes standard with the full installation). My version of git, which you can see by executing:

```
git --version
```

is 1.7.0.4.

Execute these commands:

```
git config --global user.name "Your Name"  
git config --global user.email "your@email.address"  
git config --global core.editor vim
```

They will create a file named .gitconfig in your home directory storing the parameters that you have entered (in some form which you can easily inspect and later edit by hand). Your specified identity will be used to automatically annotate your changes.¹

¹ Note that all the modifications introduced until June 28 (which have been automatically carried over from the CVS repository) have a dummy author.



The central repository

You have to log on to your account that has been created for you at olsonet.unfuddle.com and provide your public ssh key, which will give your git smooth access to the repositories. You will receive an automatically generated e-mail with your user ID and a temporary password. Sign in as instructed. Then, in your *Personal Settings*, move down to *Public Keys* and click on *New Public Key*. Enter a name for the key (e.g., My New Laptop) and paste your public ssh key into the *Value* field (do I have to explain this?). This will allow you to access the repositories. If you think you need a write (git push) access (and you don't have it), please let me know.

Having entered the new key, and saved the changes, click on the *Repositories* tab. You will see the three repositories along with their URLs. Those URLs will allow you to connect to them from your machine. Note that unfuddle offers some features (like web browsing through the repositories). Please indulge. And if you discover something valuable, please let me know. So far we have been using the very basics.

Back to the local system

Fetch the stuff to your local machine. Move to the directory where you want to keep it and execute:

```
git clone git@olsonet.unfuddle.com:olsonet/side.git SIDE
git clone git@olsonet.unfuddle.com:olsonet/picos.git PICOS
git clone git@olsonet.unfuddle.com:olsonet/vuee.git VUEE
```

In fact, the last arguments can be full directory paths, so you can execute these commands from anywhere. They will fetch the repositories and checkout the *master* contents of each of them.

Logging changes

Before June 28, we used to have a single file in the root directory of a package, e.g., RTAGS in PICOS, where we would put tags (typically something like R100811A) specifying the date and followed by a description of what has changed. After the switchover to git, we agreed to follow this approach:

1. All commits receive brief descriptions (git prompts for them automatically). Such a description is typically one or two sentences long and refers to the most significant change introduced in the commit.
2. Not all commits must be tagged (in the old sense), but if a commit is tagged, then the tag should appear in the commit's description.
3. Every developer has their RTAGS file in the root directory of the respective package, e.g., RTAGS_PG, RTAGS_WO. This is where we put more elaborate descriptions of the changes. The tag format is *xyymmddda*, where *xx* are the initials of the developer, *yymmdd* is the date (e.g., 100729), and *a* is a letter (typically A) for the occasion of multiple tags introduced on the same day.
4. While not all commits are tagged, all tags appearing in RTAGS_... correspond to commits.

Here is a sample entry from RTAGS_PG in PICOS:



PG100730A:

A few cleanups in the CC1100 driver mostly related to option constants. The options are described in Docs/PICOS/cc1100.odt.

and here are two sample commit messages:

PG100811A: a few tweaks to picomp (read RTAGS_PG).

Not tagged: a channel change command added to VUEE/RFTEST (see page 9 of the doc).

The original tag files will not be updated any more (but they have been retained). As the "official" documentation lags behind a bit (to put it mildly – sorry, sorry ... I will try to catch up in the next few days), the tag files (RTAGS...) are the most up-to-date source of wisdom.

A git tutorial of sorts

For a quick illustration, move to PICOS and do this:

```
git checkout GENESIS
```

After two or three seconds you are inside the GENESIS release. Then do:

```
git checkout master
```

and you are back in the most recent commit of the "main trunk" (using the now obsolete terminology of CVS).

Let us go through a simple modification cycle. Suppose that you are in master, and you have modified some files. When you execute:

```
git status
```

git will show you the files that have been modified and not yet *staged* for addition (this terminology is explained in chapter 2 of the book I mention above) as well as any files (like binaries, for example) that have been created anywhere in the tree, but are not being *tracked*, i.e., they do not belong to the repository.

You can *add* each of those modified files (the same applies to any new files to be included in the repository) with:

```
git add filename
```

or

```
git add -A
```

if all of them are to be tracked (status doesn't report any files that should not be in the repository). Then you can do:

```
git commit
```

```
git tag tagname (e.g., R100629A)
```



Note that up to this point all the changes, including the commit, apply to your copy of the repository. To send them to the central repository, do this:

```
git push --all
git push --tags
```

The second operation is not needed if you haven't introduced any tags that should be copied to the central repository. This pretty much covers the way we used to work with CVS. Here is the transcript of a somewhat more sophisticated way to introduce and test changes:

```
git checkout master
git branch mytest
git checkout mytest
```

The last two commands can be compressed into a single:

```
git checkout -b mytest
```

The sequence creates a new branch which starts being identical to master. You can now edit/create files, run tests, and so on, without affecting anything at master (to which you can always easily revert). You can also commit the changes (they will be committed to mytest); you can even push them to the central repository (e.g., to show to other people). When you are ready, you do:

```
git checkout master
git merge mytest
```

and the changes will be merged (*fast forwarded*) into master. Let us go through a complete exercise using specific files, branches, and tags. This will pretty much cover all I have ever been doing with CVS (in fact quite a bit more ;-). So we start:

```
cd PICOS
git checkout master
git checkout -b mytest
```

Our objective is to introduce some changes, pretend to test them (keeping master intact for as long as possible), then merge the new changes into master and commit them (tagging the commit in our traditional way).

While operating within mytest, we move to Apps/VUEE/survey and edit two files: app_peg.cc and app_diag_peg.cc (changing the copyright year to 2010). We compile the praxis for a real node:

```
mkmk WARSAW
make
```

Then we move back to PICOS and add a note at the end of RTAGS. (tagged with R100627A). Now, when we say:

```
git status
```

we get this:

```
# On branch mytest
# Changed but not updated:
```



```
# (use "git add <file>..." to update what will be
# committed)
# (use "git checkout -- <file>..." to discard changes in
# working directory)
#
#       modified:   Apps/VUEE/survey/app_diag_peg.cc
#       modified:   Apps/VUEE/survey/app_peg.cc
#       modified:   RTAGS
#
# Untracked files:
# (use "git add <file>..." to include in what will be
# committed)
#
#       Apps/VUEE/survey/.gdbinit
#       Apps/VUEE/survey/Image_peg
#       Apps/VUEE/survey/Image_peg.a43
#       Apps/VUEE/survey/KTMP/
#       Apps/VUEE/survey/Makefile
#       Apps/VUEE/survey/gdb.ini
no changes added to commit (use "git add" and/or
"git commit -a")
```

Note that the untracked files result from the compilation (they do not belong to the repository). We can ignore them and add the relevant files by hand, e.g.,

```
git add Apps/VUEE/survey/app_diag_peg.cc \
      Apps/VUEE/survey/app_peg.cc RTAGS
```

(this is supposed to be a single line), or remove the untracked files:

```
cd Apps
./cleanup
```

do git status again (for certainty), and say:

```
git add -A
```

to add all the modified files.

A digression: you do not absolutely have to execute git at the root level of the package (where the special .git directory is kept). Apparently, git searches for .git upwards from where you call it. In such a case, the paths of all listed files (like in the above output) will be relative to where you are (so they always can be sensibly pasted as arguments of commands).

The next step is to commit the changes, i.e.,

```
git commit
```

This will open the text editor (the one you configured in as your favorite) asking you to enter a commit message.

So let us enter this:

```
PG100627A: Dummy changes testing the new git setup.
```



and exit the text editor. This completes the commit. Now we can assign the tag to it:

```
git tag PG100627A
```

The changes (including the commit) have gone into branch mytest. This branch can still be worked on (more commits) without affecting master (we can switch back and forth between master and mytest). When we are ready, we do:

```
git checkout master
git merge mytest
```

which merges mytest into master. In this case the operation "fast forwards" master to our new commit, because mytest descends directly from master (which doesn't have to hold generally). Now, we can remove the temporary branch (which at this time points to exactly the same place as master):

```
git branch -d mytest
```

and we are ready to synchronize the central repository (push the changes into it):

```
git push --all
```

This should be followed by:

```
git push --tags
```

because we have added a new tag to the local repository (and we want it to show up in the remote one). For some reason, tags have to be pushed separately. This is probably because they are not considered as useful as we have made them for us.

Whenever you want to synchronize your personal repository with the central one, you do:

```
git pull
```

The way this operation is performed is quite educational. Your repository maintains a number of pointers (heads) to various branches of the tree. One of them, master, intentionally points to the end of the "main trunk". The above command accepts optional arguments, called *refspecs*, of the form <src>:<dst>, where <src> describes a source head (in the remote repository), and <dst> is interpreted locally. By default (no explicit refsspecs), the set of effective resfspecs is described by a pattern in your .git/config file, i.e., in this section:

```
[remote "origin"]
    fetch = +refs/heads/*:refs/remotes/origin/*
    url = git@olsonet.unfuddle.com:olsonet/picos.git
```

which describes the remote repository associated (by default) with your local repository. The *fetch* pattern says this: for all heads present in the remote repository, copy them to the corresponding remotes/origin/... heads in the local repository. This is what would be accomplished by a default (argument-less) git fetch command.

Note that you can reference from your local repository heads like remotes/origin/master, for example, you can do:

```
git checkout remotes/origin/master
```



which will get you to master as it looked in the remote repository when you last pulled (or fetched) from it. View this as a backup version of master (and this also applies to other branches that you fetched from the remote repository) which you are not supposed to do any development on. These are the branches to fetched into from the remote repository. Needless to say, you can branch them off and do with that stuff whatever you please.

The extra job performed by pull, in addition to fetching the "remotes", is to merge them into your refs/heads, thus, bringing your development heads up to date with the remote repository.

